

AI

What is Intelligence?

- The capacity to acquire and apply knowledge.
- The faculty of thought and reason.

Artificial Intelligence:

-Studies how to achieve intelligent behavior through computational means

The Turing Test:

- A human interrogator. Communicates with a hidden subject that is either a computer system or a human.

If the human interrogator cannot reliably decide whether or not the subject is a computer, the computer is said to have passed the Turing test.

In general, there are various reasons why trying to mimic humans might not be the best approach to AI:

- Computers and Humans have a very different architecture with quite different abilities.
- Numerical computations
- Visual and sensory processing
- Massively and slow parallel vs. fast serial

Neuroscience has been very influential in some areas of AI. For example, in robotic sensing, vision processing

Humans might not be best comparison?

- Don't always make the best decisions
- Computer intelligence can aid in our decision making

Rationality: Typically, this is a precise formal notion of what it means to do the right thing in any particular circumstance. Provides:

- A precise mechanism for analyzing and understanding the properties of this ideal behavior we are trying to achieve.
- A precise benchmark against which we can measure the behavior the systems we build.

AI tries to understand and model intelligence as a computational process.

Thus, we try to construct systems whose computation achieves or approximates the desired notion of rationality

Subareas of AI:

- Perception: vision, speech understanding, etc.
- Machine Learning, Neural networks
- Robotics
- Natural language processing
- Reasoning and decision making
- Many of the popular recent applications of AI in industry have been based on Machine Learning, e.g., voice recognition systems on your cell phone.
- Probabilistic graphical models are fundamental in machine learning.

AI Success:

- **Games:** chess, checkers, poker, bridge, backgammon... – Search
 - **Physical skills:** driving a car, flying a plane or helicopter, vacuuming...
– Sensing, machine learning, planning, search, probabilistic reasoning
 - **Language:** machine translation, speech recognition, character recognition,..
– Knowledge representation, machine learning, probabilistic reasoning
 - **Vision:** face recognition, face detection, digital photographic
 - **Commerce and industry:** page rank for searching, fraud detection, trading on financial markets...
– Search, machine learning, probabilistic reasoning
- Formalisms and algorithmic ideas have been identified as being useful in the construction of these “intelligent” systems.

Search:

- 1-One of the most basic techniques in AI
- 2-Can solve many problems that humans are not good at (achieving super-human performance)
- 3-Very useful as a general algorithmic technique for solving many non-AI problems.

Search is a computational method for capturing a particular version of this kind of reasoning.

why search:

1- Successful:

- Success in game playing programs based on search.
- Many other AI problems can be successfully solved by search.

2- Practical:

- Many problems don't have specific algorithms for solving them. Casting as search problems is often the easiest way of solving them.
- Search can also be useful in approximation (e.g., local search in optimization problems).

Some critical aspects of intelligent behavior, e.g., planning, can be naturally cast as search.

Limitations of Search:

- Search only shows how to solve the problem once we have it correctly formulated.

formulate a problem as a search problem:

1-state space

2-actions

3-initial state and goal

4-heuristics

Inputs for search algorithms:

1-initial state

2-successors

3-goal test

4-actions cost

Outputs:

- a sequence of states leading from the initial state to a state satisfying the goal test.

Template Search Algorithms:

- The search space consists of **states** and actions that move between states.
- A **path** in the search space is a **sequence** of states connected by actions.
- The search algorithms perform search by examining alternate paths of the search space. The objects used in the algorithm are called nodes—each node contains a path.

We maintain a set of Frontier nodes also called the OPEN set.

These nodes are paths in the search space that all start at the initial state.

Selection Rule:

The order paths are selected from OPEN has a critical effect on the operation of the search:

- Whether or not a solution is found
- The cost of the solution found.
- The time and space required by the search.

All search techniques keep OPEN as an ordered set (e.g., a priority queue).

Critical properties of Search:

- 1- **Completeness** (will it find the solution if it exists)
- 2- **Optimality** (Will it always find the least cost solution)
- 3- **Time complexity** (what is the max number of paths that can be generated)
- 4- **Space Complexity** (what is the max number of paths that have to be stored)

Uninformed search techniques:

- 1-Breadth-First. 2-Depth-First. 3-Uniform-Cost 4-Depth-Limited
- 5-Iterative-Deepening search

Breadth-First: Place the new paths that extend the current path at the end of OPEN.

Completeness:

- 1- The length of the path removed from OPEN is non-decreasing
- 2- **All** shorter paths are expanded prior before any longer path.
- 3- Hence, eventually we must examine all paths of length d , and thus find a solution if one exists

Optimality:

- 1- will find shortest length solution
- 2- shortest solution not always cheapest solution if actions have varying costs.

Time Complexity: $O(b^{d+1})$

Space Complexity: $O(b^{d+1})$

Space complexity is a real problem.

Typically run out of space before we run out of time in most applications.

Uniform-Cost Search: Keep OPEN ordered by increasing cost of the path.

- Identical to Breadth first if each action has the same cost.

Completeness:

- 1- If each transition has costs $\geq \epsilon > 0$.

Optimality:

- 1- Finds optimal solution if each transition has cost $\geq \epsilon > 0$.
 - 2- Explores paths in the search space in increasing order of cost.
 - 3- So must find minimum cost path to a goal before finding any higher costs paths.
- Time Complexity:** $O(b^{(C^*/\epsilon)})$ where C^* is the cost of the optimal solution.
- Space Complexity:** $O(b^{(C^*/\epsilon)})$
- when a path to a goal state is expanded the path must be optimal (lowest cost). There may be many paths with cost $\leq C^*$: there can be as many as bd paths of length d in the worst case.

Depth-First Search: Place the new paths that extend the current path at the front of OPEN.

Completeness:

- 1- Infinite paths? Cause incompleteness!
- 2- Prune paths with cycles (duplicate states) We get completeness if state space is finite

Optimality: NO!

Time Complexity: $O(b^m)$ where m is the length of the longest path

Very bad if m is much larger than d (shortest path to a goal state)

Space Complexity: $O(bm)$, linear space

A significant advantage of DFS

Depth-Limited Search:

- Breadth first has space problems, Depth first can run off down a very long (or infinite) path.
- Perform depth first search but only to a pre-specified depth limit D .
- THE ROOT is at DEPTH 0, ROOT is a path of length 1.
- No node representing a path of length more than $D+1$ is placed on OPEN.
- We “truncate” the search by looking only at paths of length $D+1$ or less.
- Now infinite length paths are not a problem.
- But will only find a solution if a solution of $DEPTH \leq D$ exists.

Iterative-Deepening Search:

-Solves the problems of depth-first and breadth-first by extending depth limited search.

-Starting at depth limit $L = 0$, we iteratively increase the depth limit, performing a depth limited search for each depth limit.

-Stop if a solution is found, or if the depth limited search failed without cutting off any nodes because of the depth limit.

- If no nodes were cut off, the search examined all paths in the state space and found no solution then no solution exists.

Completeness:

-Yes, if a minimal depth solution of depth d exists.

Time Complexity: $O(b^d)$ Most nodes lie on bottom layer.

-BFS can explore more states than IDS!

-In fact, IDS can be more efficient than breadth first search: nodes at limit are not expanded. BFS must expand all nodes until it expands a goal node. So the bottom layer it will add many nodes to OPEN before finding the goal node.

Space Complexity: $O(bd)$

-Will find shortest length solution which is optimal if costs are uniform.

- If costs are not uniform, we can use a “cost” bound instead.

- Only expand paths of cost less than the cost bound.

- Keep track of the minimum cost unexpanded path in each depth first iteration, increase the cost bound to this on the next iteration.

-This can be more expensive. Need as many iterations of the search as there are distinct path costs.

Path checking:

-Ensure that the state c is not equal to the state reached by any ancestor of c along this path.

-Paths are checked in isolation!

Cycle Checking:

-Keep track of all states previously expanded during the search.

-When we expand n_k to obtain child c , ensure that c is not equal to **any** previously expanded state.

-This is called cycle checking, or multiple path checking.

-Higher space complexity (equal to the space complexity of breadth-first search).

Heuristic Search(informed search):

In uninformed search, we don't try to evaluate which of the nodes on OPEN are most promising. We never “look-ahead” to the goal.

- The idea is to develop a domain specific heuristic function $h(n)$.
- $h(n)$ guesses the cost of getting to the goal from node n (i.e., from the terminal state of the path represented by n).

heuristics are domain specific.

- If $h(n_1) < h(n_2)$ this means that we guess that it is cheaper to get to the goal from n_1 than from n_2 .

Greedy best-first search:

- We use $h(n)$ to rank the nodes on OPEN
- Always expand node with lowest h -value.
- We are greedily trying to achieve a low-cost solution.
- However, this method ignores the cost of getting to n , so it can lead astray exploring nodes that cost a lot but seem to be close to the goal.

A* search:

- Take into account the cost of getting to the node as well as our estimate of the cost of getting to the goal from the node.

Define an evaluation function $f(n)$

$$f(n) = g(n) + h(n)$$

- $g(n)$ is the cost of the path represented by node n
- $h(n)$ is the heuristic estimate of the cost of achieving the goal from n .
- Always expand the node with lowest f -value on OPEN.
- The f -value, $f(n)$ is an estimate of the cost of getting to the goal via the node (path) n .

UNTIL PAGE 90!